

# Thoughts, Experiences and Considerations with Throughput and Latency on High Volume Stream Processing Systems Using Cloud Based Infrastructures.



# How I Bled All Over Onyx




Work for....




Wrote this....

Tweets here:  
[@jasonbelldata](https://twitter.com/jasonbelldata)




Tweet heckling  
is perfectly  
allowed.

#clojurex






Actually any  
heckling is  
perfectly  
allowed.



# One Day I Got a Question



▶ “We want to stream in 12TB of data a day..... Can you do that?”




Yeah! We're  
gonna rock  
like Cameo  
'86



# A Streaming Application... kind of.

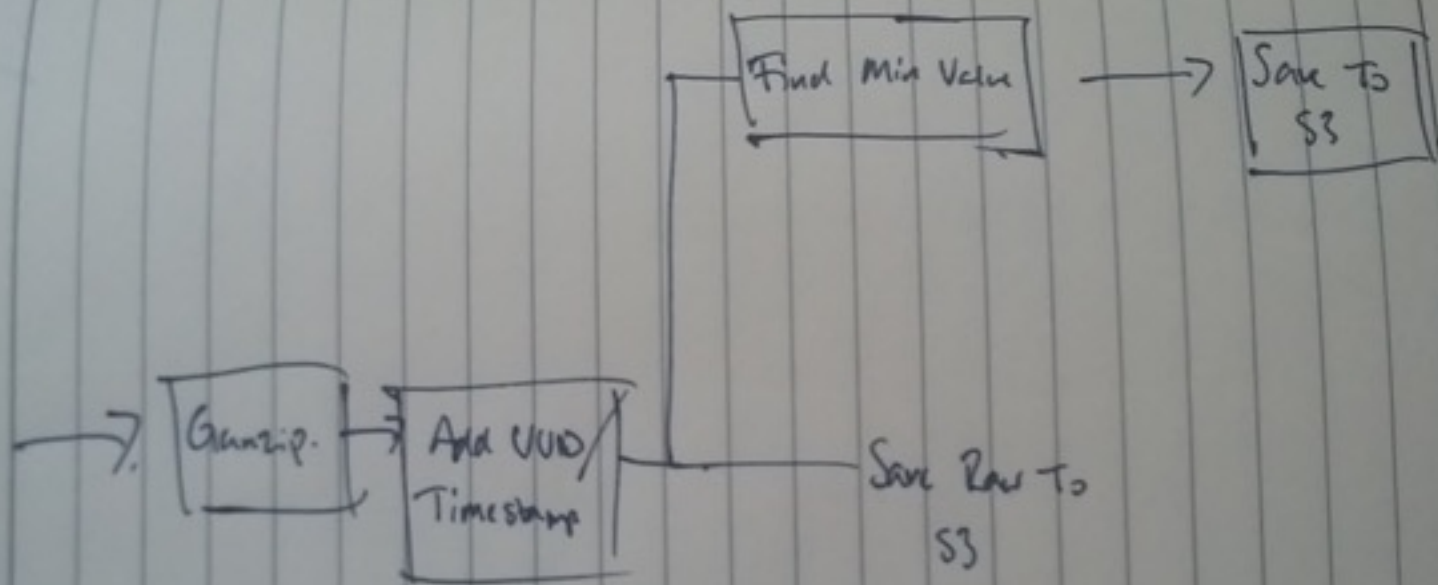
How Everyone Draws Streaming Apps.



It's never that simple.

# What is Onyx?

- ▶ Distributed Computation Platform
- ▶ Written 100% in Clojure
- ▶ Great plugin architecture (inputs and outputs)
- ▶ Uses graph Direct Acyclic Graph workflow
- ▶ I spoke about it at ClojureX 2016
- ▶ It's very good....





# Onyx: Workflows as maps.



- ▶ `[[:in :add-fields]`
- ▶ `[:add-fields :dump-row]`
- ▶ `[:add-fields :find-min]`
- ▶ `[:find-min :dump-min]]`

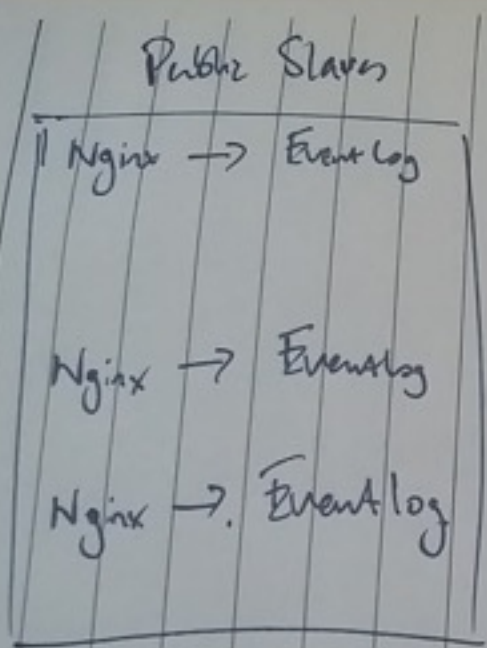
# Onyx: Basic Peer Design

- ▶ 8 Peers
  - ▶ Per Kafka Partition (3)
  - ▶ Input task (deserialisation) (1)
  - ▶ Assign Date + UUID (1)
  - ▶ Write row to S3 (1)
  - ▶ Find min value row (1)
  - ▶ Write min value to S3 (1)


# Terraform/Marathon/Mesos Deployment

- ▶ 3 Masters (m4.large)
- ▶ 3 Public Slaves (t2.medium)
- ▶ 3 Private Slaves (m4.xlarge)
- ▶ 4TB Volumes
- ▶ We'd see how it goes from there, at worst d2 instances will do.





T2 Med.




M4 Xlarge

# Great! Let's Build!




# Phase 1 – Testing at 1% Volume



- ▶ It's working!
- ▶ Stuff's going to S3!
- ▶ Tell the client the good news!



# Phase 2 – Testing at 2% Volume



- ▶ It's working!
- ▶ Stuff's going to S3!
- ▶ Tell the client the good news!

# Phase 3 – Testing at 5% Volume







This is new territory and I'm scared.  
I may have pushed Onyx 4sd to the mean.



# Know Thy Framework

- ▶ Aeron buffer size =  $3 \times (\text{batch-size} \times \text{segment size} \times \text{connections})$ 
  - ▶  $(1 \times 512\text{k} \times 8) \times 3 = 4\text{mb}$
  - ▶ Aeron default buffer is 16mb but then the **twist**
    - ▶ Onyx segment size =  $\text{aeron.term.buffer.size} / 8$
    - ▶ Max message size of 2mb

# Know Thy Framework

- ▶ Each peer has it's own lifecycle
  - ▶ If one task dies the then the whole peer structure can't process data, simple.
  - ▶ Make sure you handle lifecycle exceptions.

```
(defn handle-exception [event lifecycle lifecycle-phase e]
  (t/info "Caught exception: " e)
  (t/info "Returning :restart, indicating that this
          task should restart.")
  :restart)
```

# Know Thy Framework

- ▶ Each peer has it's own lifecycle
  - ▶ Heartbeat default is 10 seconds
    - ▶ `:onyx.peer/subscriber-liveness-timeout-ms`
    - ▶ `:onyx.peer/publisher-liveness-timeout-ms`
    - ▶ **If the heartbeat times out the whole job shutdowns.**
    - ▶ So I changed it to 30 seconds.



# Onyx in Docker containers.


- ▶ `--shm-size` ended up being 10GB
  - ▶ OOM killers were frequent
  - ▶ Java logging is helpful but still hard to deal with.

```
CGROUPS_MEM=$(cat /sys/fs/cgroup/memory/memory.limit_in_bytes)
MEMINFO_MEM=$((($ (awk '/MemTotal/ {print $2}' /proc/
meminfo)*1024)) MEM=$((($MEMINFO_MEM>$CGROUPS_MEM?$CGROUPS_MEM:
$MEMINFO_MEM)) JVM_PEER_HEAP_RATIO=${JVM_PEER_HEAP_RATIO:-0.6}
XMX=$((awk '{printf("%d", $1*$2/1024^2)}' <<< " ${MEM} $
{JVM_PEER_HEAP_RATIO} ") # Use the container memory limit to
set max heap size so that the GC # knows to collect before
it's hard-stopped by the container environment, # causing OOM
exception.
```

# One Friday evening...



...two things happened...





1. Watched “12 Days of Christine”  
and cried.



## 2. I did a rewrite in Kafka Streams.


- ▶ Moved the pure Clojure functions into the streams architecture.
- ▶ Used Amazonica to write to AWS S3.

# Kafka Streams rewrite...

- ▶ Took 2 hours to write and deploy to DCOS.



# Kafka Streams rewrite...



- ▶ Took 2 hours to write and deploy to DCOS.
- ▶ And, by this stage, was slightly tipsy.

# Kafka Streams rewrite...



- ▶ In deployment Kafka Streams didn't touch more than 2% of the CPU/Memory load on the instance.
- ▶ Max memory was 790Mb compared to 8Gb Onyx jobs.

# Monday Morning...

- ▶ No restarts, still processing data from all partitions!




# Monday Morning...

- ▶ Now might be a good time to tell the CTO. ;)

# So why did Onyx bleed so bad?

- ▶ Unzipping messages during deserialisation introduces unknowns across the peers. Very hard to monitor and predict.
- ▶ Onyx isn't designed for large messages (kb not mb)



```
(defn deserialize-gzip-message [bytes]
  (try (-> bytes
           ByteArrayInputStream.
           GZIPInputStream.
           io/reader
           slurp)
        (catch Exception e {:error e})))
```



# Am I saying Onyx is bad?

- ▶ Definitely **not**. It's wonderful. You need to confirm it's wonderful for your use case.

# Key Takeaways

- ▶ Design on paper first.
- ▶ Think about every function the message will visit.
- ▶ Think about the fail safes on each function/task (size, timeouts, rogue messages).

# Key Takeaways

- ▶ Know the key Kafka topic settings.
- ▶ LOG.RETENTION.HOURS,  
LOG.RETENTION.MINUTES,  
LOG.RETENTION.MS,  
LOG.RETENTION.BYTES,  
MESSAGE.MAX.BYTES
- ▶ <https://dataissexy.wordpress.com/2017/03/10/kafka-diaries-topic-level-settings-you-cant-ignore-part-1-data-streaming/>



# Key Takeaways

- ▶ Every task in the workflow will incur latency and adds up over the duration of the stream.
- ▶ Work out your TTL backstop, log retention, on Kafka (default 168 hours).

# Key Takeaways

- ▶ Using Zookeeper? Monitor it with your life (or someone else's)
- ▶ Reduce disk write i/o on Kafka, throw as much RAM at the brokers as you can.
- ▶ On Kinesis you don't get the luxury of tuning and is about 70% slower on throughput.

# Key Takeaways

- ▶ If your customer is not on a 24/7 SLA then work out how much data is going to pass through on the down days, then plan storage accordingly. **Kafka will log while consumers are dead.**



# Key Takeaways


- ▶ Think the Streaming Rule of 72

# The Streaming Rule of 72



- ▶ Work out % gain
  - ▶  $(130 \text{ msg/s} - 100 \text{ msg/s}) / 100 = 0.3$   
\* 100 = 30%
- ▶  $72 / 30 = 2.4$
- ▶ You will double the message throughput every 2.4 seconds.

# Thank you.



- ▶ Ask me questions anytime either here or via Twitter, LinkedIn, email and so on.
- ▶ If I can help I will help.